

RYO-CHAN Virtual Hackathon 2026 — Project Submission Form

Why this file exists. The organiser's official form is linked from GET `https://app-ryochan.com/api/hackathon/config` as `https://ryobuild.com/project-submission-form.pdf`, but that URL does not serve a PDF. Checked again on 2026-09-06: HTTP 200, content-type: text/html, 2 034 bytes of the site's SPA shell (same for `MCP-Builder-Guide.pdf`; only the `.md` guide is a real file). The Help Desk has been asked for a working link. Until it arrives this form reproduces every field of the platform's own `HackathonSubmissionFields` schema (PUT `/api/hackathon/submission/draft`), so the committed answers match what will be submitted through the API.

field	value
team_name	RYO Arena
participant_name	Pugar Huda Mantoro
email	hudapugar@gmail.com
discord_id	Lynx (hajislamet)
github_username	PugarHuda
project_name	RYO Arena
tracks	track_1, track_2, track_3
repo_url	pending — organiser's private repo not yet issued; working mirror: <code>https://github.com/PugarHuda/ryo-arena</code> (private)
demo_video_url	pending upload
x_post_url	pending post
agree_rules	yes
confirm_no_secrets	yes — <code>.env</code> is gitignored, <code>.env.example</code> carries names only, <code>git grep</code> for <code>ryo_mcp_</code> , <code>tvly-</code> , <code>sk-or-v1-</code> and <code>VENICE_INFERENCE_KEY_</code> returns only placeholders and test doubles (verified 2026-09-06)

project_description

RYO Arena makes an AI trading opinion auditable. Every decision is a receipt you can re-run: `arena replay <id>` rebuilds it from the stored evidence and cached model output and prints identical: true, so nothing can be rewritten after the fact. Each cited number carries its dotted RYO path and is read back out of the evidence, never retyped by the model; citations pointing at absent evidence are dropped in code; null is never turned into 0. Independent sources audit RYO itself - exchange medians check its price, Wilder RSI/ATR recomputed from public OHLC check its indicators - and the judge refuses to size a trade when they disagree. Three specialists (macro, technician, narrative) debate all six builder tools, weighted by their own Brier score once calls resolve. Track 3: four skills on RYO's own `/api/skills` paths, in

RYO's envelope. Track 2: a dashboard that diffs each receipt against the last, ranked by impact. Read-only; no orders.

Evidence the judges can check without any key

- Hosted dashboard: <https://ryo-arena.vercel.app> (read-only ledger snapshot, `/api/health` reports `ryo_key_set: false` — nothing is disguised as live).
- `ARENA_DB=data/demo.db uv run arena serve, uv run arena skill run price_crosscheck '{"symbol":"SOL"}', uv run pytest -q` (110 tests, no network).

Declarations

- Read-only research tool. It never places, signs or routes an order; positions are paper only.
- No RYO builder key has ever been committed. Fixtures under `tests/fixtures/` are hand-built and labelled; nothing simulated is ever presented as live (`data_mode` is carried through and the risk layer refuses to size on `simulated`).
- Third-party libraries are disclosed in README under "Disclosed third-party libraries". All application code was written during the hackathon.